



Infosys Challenge Tracks for HackMTY 2026

2. The Forensic Auditor

Section	Content
About your company	Infosys is a global technology and consulting company with deep experience in finance, accounting, and risk. From its Monterrey center it has done bookkeeping, compliance, and fraud-investigation work for clients across the Americas since 2007.
The problem today	Invoice fraud in Mexico is not a neat list of flagged rows. It is a scheme hidden inside a real company's books: a fake supplier billing for work never done, a kickback routed through a shell company, or sales faked to inflate the numbers. Mexico's tax authority, SAT, publishes a blacklist of fake-invoice companies under Article 69-B, but by the time a supplier appears on it, the company has already claimed the deductions and is on the hook. Today's tools flag odd items one at a time and leave a person to connect them. Nobody follows the money all the way, nobody builds the proof, and honest suppliers who simply look odd get accused too.
The big challenge	Given a company's books and only the hint that something is wrong, can an AI agent find the fraud, follow the money, and prove it, without accusing anyone it cannot back up?
Possible approaches	(1) Investigate step by step: form a theory, search the ledger, invoices, and bank records, follow a lead, and change course when it dead-ends. (2) Use simple detectors (blacklisted suppliers, payments that do not match invoices, money that moves in a circle) to point the agent at what is worth digging into. (3) Build an evidence trail for every accusation and refuse to name a supplier it cannot back with a clear rule broken and a peso amount.
Resources for hackers	Teams build on real, free resources that exist online. SAT publishes the official Article 69-B list of companies that issue fake invoices (EFOS), a real, downloadable Mexican dataset, along with the CFDI 4.0 invoice schemas. IBM AMLSim (open-source) generates the money-flow rings and shell-company patterns behind kickbacks and round-tripping. Public financial-fraud datasets, such as the IEEE-CIS set on Kaggle, give baselines for the anomaly checks. Teams assemble a company data estate from these pieces, a ledger, invoices, bank records, and a supplier list, then build the investigation loop. Note on tools: this track makes many AI calls per investigation, so a local model (Ollama) with caching is safer than the Gemini free tier alone, which can hit daily limits.
What to build	A forensic agent that investigates records it has never seen and hands in a case file: the scheme, the suppliers involved, the evidence trail, and the peso amount, plus a short list of leads it chose not to chase and why. Deliver working code and a 3-minute live demo where judges hide a fresh scheme in the data, the agent traces the money on screen, then answers one surprise question about its reasoning.
Why it matters	Every peso lost to fake invoices is taken from an honest business and the public purse, and every supplier wrongly accused loses a customer for no reason. An agent that investigates and proves, instead of just flagging, is the difference between a report that falls apart and a case a company can act on. This is the forensic work Infosys does at scale, and proving before accusing is what separates a real auditor from a guesser.
Judging criteria	Results: on records it has never seen, how much hidden fraud does the agent find and correctly prove? Judgment: does it refuse to accuse suppliers it cannot back up, and can it defend a finding when a judge asks? Feasibility: could a real finance or audit team trust and use this? Clarity: is the case file easy to follow, with a clear money trail?

3. The Courier

Section	Content
About your company	Infosys is a global technology company that builds logistics and last-mile delivery platforms for retail and e-commerce clients. From its Monterrey center it works on routing and operations across the Americas.
The problem today	Thousands of students and young people in Monterrey earn money driving for DiDi, Rappi, and Uber. The app shows a stream of orders, and the courier has seconds to accept or skip. Take the wrong order and you burn gas and time crossing the city in 40-degree heat for a few pesos. Skip too much and you earn nothing. During surge and rain the math changes every minute, and no app tells the courier the smart move, it just shows the next ping.
The big challenge	Given a stream of delivery offers, traffic, and a limited shift, can an AI agent earn the most for a courier, the way a smart, experienced driver would?
Possible approaches	(1) Decide which orders to accept and which to skip based on pay, distance, and time. (2) Batch nearby drop-offs and plan an efficient route through real Monterrey streets. (3) Learn the day's rhythm, where and when the good orders appear, and position for surge while avoiding slow or unsafe areas.
Resources for hackers	There is no ready-made simulator for this. Teams build on real, free resources that exist online. OpenStreetMap with a routing library (OSMnx or OSRM) gives the real Monterrey street network and travel times. Google OR-Tools (open-source) provides strong routing and batching baselines. Public delivery and vehicle-routing datasets, such as the Solomon VRPTW benchmark and food-delivery sets on Kaggle, supply realistic order patterns. Teams build the shift itself: the order stream, surge, and the earnings and time budget. Free AI tools (Gemma on Device, the Gemini free tier or a local model like Ollama) can be used.
What to build	An agent that plays a courier over a simulated shift, accepting or skipping orders, batching, and routing to earn the most while staying within time and keeping safe. Deliver working code and a live demo where two agents run the same fresh shift side by side, an earnings counter ticking, and a surge or a road closure hits mid-shift that they must react to.
Why it matters	Gig work is how a huge number of young people in Mexico make money, and the platforms optimize for themselves, not the driver. An agent that plays the shift well shows what fair, driver-side optimization looks like, and it is the same routing and last-mile problem Infosys solves for delivery clients.
Judging criteria	Results: on a fresh shift it has not seen, how much does the agent earn compared to a simple baseline? Judgment: does it make safe, sensible calls under surge and traffic, and can it explain a choice to a judge? Feasibility: could this realistically help a real courier during a shift? Clarity: is it easy to follow what the agent decided and why?